






Design Credit Project:

UAV Control for LiDAR with RL

Problem Statement



We implemented LiDAR and RL to the classic pursuit problem. The goal here is to reach the stationary target while avoiding all the obstacles and throughout this process the pursuer moves with constant velocity.

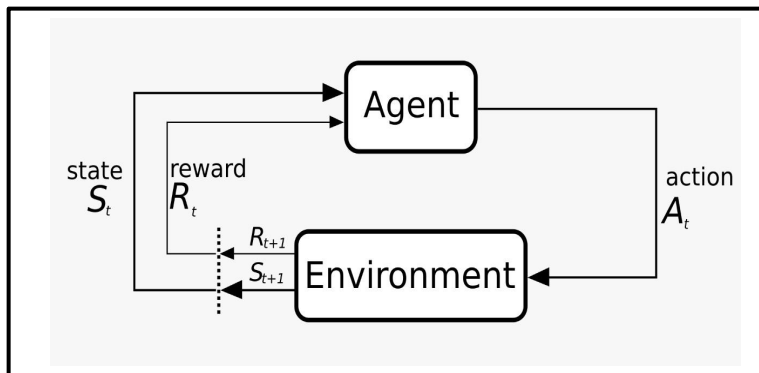


01

RL & DQN



Reinforcement learning



Theory


This is a model free,
off-policy Reinforcement
learning algorithm

$$Q^\pi(s, a) = \sum_{s' \in \mathcal{S}} P(s'|s, a) [R(s, a, s') + \gamma V^\pi(s')] \quad [1]$$

$$Q_*(s, a) = \max_{\pi} Q_{\pi}(s, a) \quad [2]$$



DQN



Working -

The weights (Q values) are updated with every step through Q update equation -

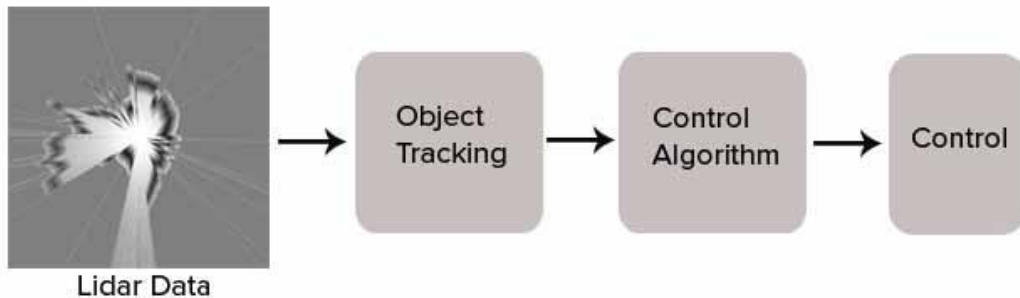
$$Q^{new}(s_t, a_t) \leftarrow \underbrace{(1 - \alpha) \cdot Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} \right)$$

learned value




Why Deep Q Networks?

Algorithmic approach



DQN approach





02

LiDAR




LiDAR

Theory

LiDAR measures the return time of laser after its reflection to determine the difference in distances of various objects.

Our LiDAR model -





03

Training & Results



Training Method

Episode Batching

- Initialize a Batch of **random** environments.
- Predict and fit every step
- Utilises parallel computation on GPU

5%	4		920/20000	[02:05<46:08,	6.89step/s]
5%	4		941/20000	[02:08<41:04,	7.73step/s]
5%	4		945/20000	[02:08<42:04,	7.55step/s]
5%	5		1009/20000	[02:17<40:34,	7.80step/s]
5%	5		1015/20000	[02:17<43:19,	7.30step/s]
6%	6		1203/20000	[02:43<42:00,	7.46step/s]
6%	6		1246/20000	[02:48<41:54,	7.46step/s]

Naive step training

- Only initialize a single random environment
- Predict every step, fit after sufficient data is learned
- Complex calculations that run on the CPU

5%	5		1095/20000	[1:40:30<24:32:52,	4.67s/episodes]
5%	5		1099/20000	[1:40:48<23:24:50,	4.46s/episodes]
6%	5		1109/20000	[1:41:45<23:41:25,	4.51s/episodes]
6%	5		1143/20000	[1:45:23<30:31:27,	5.83s/episodes]
6%	5		1163/20000	[1:47:16<25:31:40,	4.88s/episodes]
6%	6		1242/20000	[1:55:41<28:45:27,	5.52s/episodes]
6%	6		1279/20000	[1:59:24<30:04:12,	5.78s/episodes]

Various Trained Models-

Three model based on different observations given to the pursuer -

Model .a

Instantaneous
positions

Model .b


Instantaneous
positions with
pursuer's velocity
vector

Model .c

Relative positions
with velocity
vector




Model 1






Model 2





Loss vs No. of episodes





Thanks for Listening!

Presented by :- Aditya Gadhavi(B19EE004)
Ajay Kumar (B19EE005)